Time to the first 128 tokens

All Flash1017 changes together vs untouched installed oMLX 0.7.1rc1 · native speculative decoding enabled

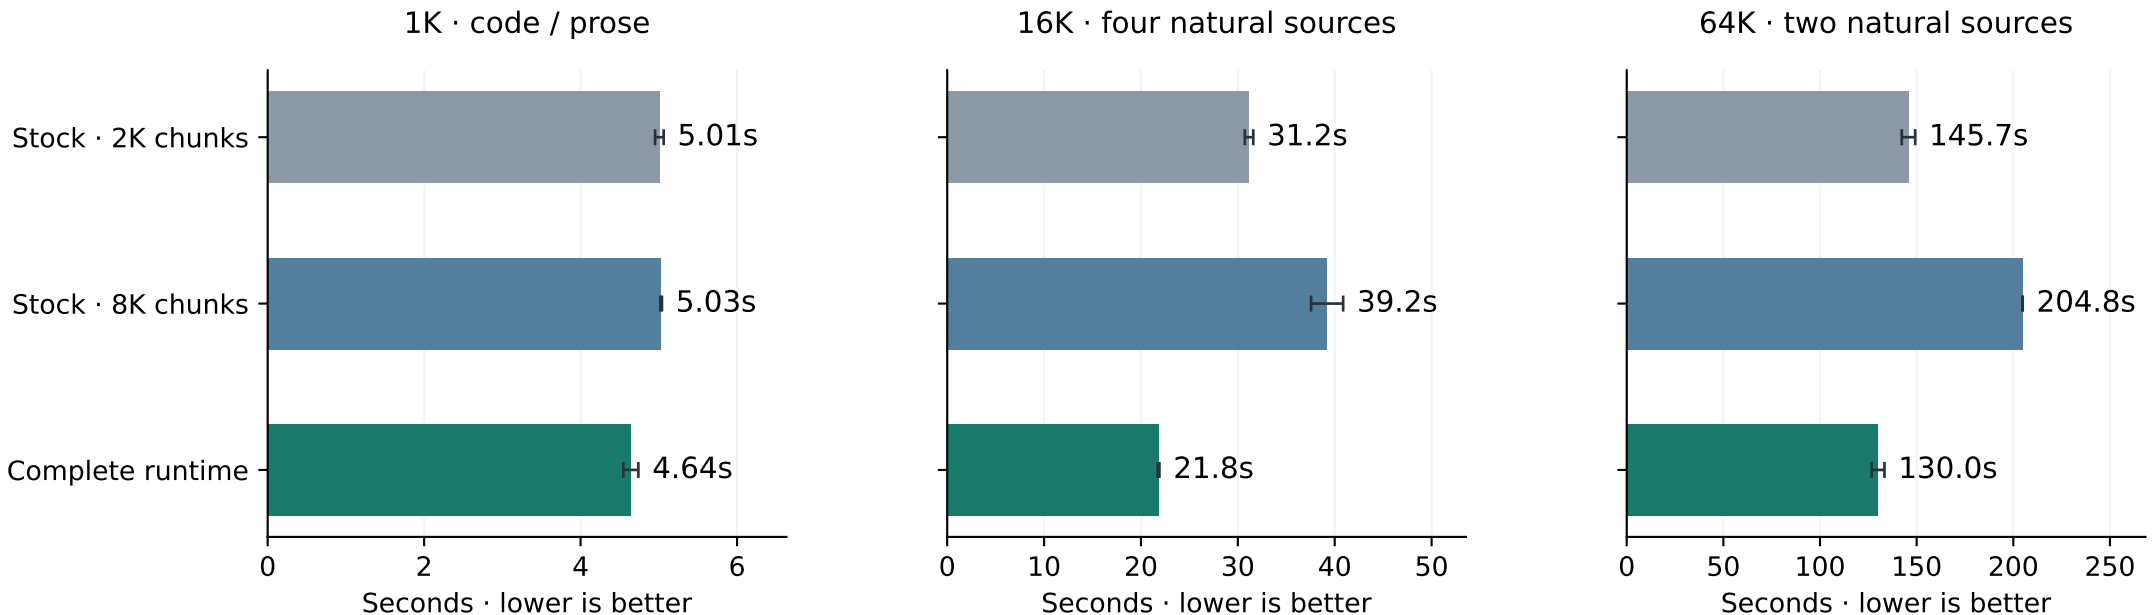

Post-hoc delivery measurement within the original 512-token runs; every prefix precedes EOS.

M5 Ultra · 80 GPU cores · 256 GiB · same oQ2.5e-mtp checkpoint · two fresh processes per configuration

Warm-model timings exclude loading and HTTP. Cache history persists as documented. Whiskers show block means, not confidence intervals.